

# Concurrency and Parallel Computing

## Paradigm

- \* Sequential Paradigm: Assume a single processor, single core architecture
- \* Concurrency vs Parallelism
  - A parallel program will use a multiplicity of independent processing units to accomplish tasks.
  - A concurrent program is one structured such that there are multiple threads of control, which gives the impression of simultaneous execution.
  - "Difference between dealing with and doing multiple things at the same time".

## Complexity

- \* Increased Design Effort
  - What is the appropriate level of division of task?  
Decomposition across many levels.
    - Blocks
    - Functions
    - Classes
    - Tasks
  - What inter-task communication is required?
    - Ideally, non-deterministic outcomes.
      - Normally, some, non-deterministic outcomes.
  - How can task correctness be maintained.
    - No longer any guarantees of order

- \* Increased Debugging Effort
  - Execution is potentially non-deterministic
  - No two executions will yield the same sequence of operations that lead to a bug
  - Creates another area in programs for bugs to hide and go unnoticed until some critical moment.
- \* Nasty Side Effects.
  - Even if programs are supposedly bug-free:
    - No compilation errors.
    - No logical errors
    - No runtime errors
  - Can have concurrency issues:
    - Race conditions
    - Dead lock
    - Starvation
  - Race condition: Two tasks that are coordinated in a sequential order is correct. Two tasks that are not coordinated and executing concurrently is wrong.
  - Dead lock: Tasks are harder to notice each other when the tasks and resources are many and possibly distributed across many different systems.
  - Resource Starvation: Task A may have higher priority, or the scheduling algorithm is unfair.

## Concurrency in GO

\* Provides simple features to investigate concurrency with minimal cognitive overhead (i.e. no additional libraries / tools to learn).

### - Goroutines

- Lightweight processes that can be created and execute function concurrently

### - go f()

### - Channels

- Communication mechanism between goroutines

```
• go func() { message <- "ping" }()
```

```
• msg := <- messages
```

## Introduction to GO

### \* Key Features

- Statically typed language

- Type inference
- Fast compilation

- Remote package management

- Using the delightful invocation go get.

- Garbage collection

- Automatic memory management

- Unrequired object's memory is reclaimed

- Composition over inheritance

- Less time worrying about complex type hierarchies

### - Build-in Concurrency

- Primitives in the language, not an extra library
- Intuitive spawning of goroutines to complete tasks.
- Model of communication based on established theory (csp)

### - Concurrency Management

- Tool support to detect concurrency bugs  
(e.g. Race conditions)

## The Inevitable Hello World

```
package main
```

```
import "fmt"
```

```
func main() {
```

```
    fmt.Println("Hello world")
```

```
}
```

## Channel creation and Communication

```
Pinger, Ponger, Printer (CPPP)
```

```
func main() {
```

```
    var c chan string = make(chan string)
```

```
    go pinger(c)
```

```
    go ponger(c)
```

```
    go printer(c)
```

```
    var input string
```

```
    fmt.Scanln(&input) }
```

```
func pinger(c chan string) {
```

```
    for i:=0; ; i++ {
```

```
        c <- "ping"
```

```
}
```

```
func ponger(c chan string) {
```

```
    for i:=0; ; i++ {
```

```
        c <- "pong"
```

```
}
```

```
func printer(c chan string) {
```

```
    for {
```

```
        msg := <-c
```

```
        fmt.Println("printer received:", msg)
```

```
        time.Sleep(time.Second * 1)
```

```
}
```

```
}
```

## Buffered Channels

### \* Unbuffered Channels:

- Channels can be created that have no buffer

- Sender goroutine will block, until a Receiver consumes the message sent into the channel

- Creation: `ic := make(chan int)`

### \* Buffered Channels

- Sender blocks only while the message has been copied into the channel
- The channel has a given capacity (size of the buffer)
- If buffer is full, sender waits for some receiver.
- Creation: `ic := make(chan int, 10)`

## Synchronisation

### \* Syncing goroutines

- Use a shared channel to coordinate when both tasks have ended
- Main(M) will wait for the Worker(W) to finish and signal over the shared channel.

### \* Channel Synchronisation

```
func worker(done chan bool) {
```

```
    fmt.Println("working...")
```

```
    time.Sleep(time.Second)
```

```
    fmt.Println("done")
```

```
    done <- true
```

```
}
```

```
func main() {
```

```
    done := make(chan bool, 1)
```

```
    go worker(done)
```

```
    <- done
```

```
}
```

## Blocking & Non-blocking Operations

### \* Using select

- Select is the switch statement for channels
  - Choose first available channel
  - Randomly choose if multiple channels ready
  - If none are ready, block or take default case.

### - In Go:

```
func main() {
    c1 := make(chan string)
    c2 := make(chan string)
    go func() {
        for {
            c1 <- "from 1"
            time.Sleep(time.Second * 2)
        }
    }()
    go func() {
        for {
            c2 <- "from 2"
            time.Sleep(time.Second * 3)
        }
    }()
    go func() {
        for {
            select {
                case msg1 := <- c1:
                    fmt.Println(msg1)
                case msg2 := <- c2:

```

```
fmt.Println(msg2)
```

```
    }
    var input string
    fmt.Scanln(&input)
}
func main() {
    c1 := make(chan string, 1)
    go func() {
        time.Sleep(time.Second * 2)
        c1 <- "result"
    }()
    select {
        case res := <- c1:
            fmt.Println(res)
        case <- time.After(time.Second * 1):
            fmt.Println("timeout!")
    }
}
```

## Using a Channel as a Queue

### \* Processing a queue

- A group of goroutines share a channel and add messages, creating a queue of work
- A single goroutine process the message in FIFO order
- Using range: iterate over each element in queue.

## Data Races

### \* Resolving Race Conditions

- A data race occurs when two or more threads access the same variable (memory) concurrently and at least one of the accesses is write.
- The solution is to synchronise access to all mutable data.
  - Using locks (to protect shared memory)
  - Using channels (to communicate activity)
- Don't communicate by sharing memory; share memory by communicating.

## Deadlock

- \* Go has runtime support for detecting a deadlock situation

## The Dining Philosophers

- \* Fun question.