

Parallel Computations for Large-Scale Problems

Lecture 2. Introduction

Def: T_s^* → execution time of the fastest serial algorithm.

T_p → total execution time of P processors

Parallel speedup $S_p = T_s^* / T_p$.

Efficiency in distributed memory machines

- * Data Parallelism : similar processing
- * Data Partitioning : partition into adjacent regions
- * Relaxed Algorithm : Embarrassingly parallel : No sync
- * Synchronous Iterations : Do synchronous.
- * Replicated Workers : Central pool.
- * Pipelined Computation : arranged in a structure

Lecture 2. Performance Evaluation

Amdahl's Law : Assume serial part $t_s = f T_1 \Rightarrow$

$$T_p = f T_1 + \frac{(1-f) T_1}{P} \Rightarrow S_p = \frac{T_1}{T_p} = \frac{P}{1+(P-1)f}$$

$$\Rightarrow \lim_{P \rightarrow \infty} S_p = \frac{1}{f}$$

Gustafson's Law : Assume T_p constant. \Rightarrow

$$T_p = f' T_p + (1-f') P T_p \Rightarrow S_p' = f' + (1-f') P$$

Def: Parallel efficiency $\eta_p = S_p / P$.

Cost $C = T_p P$.

Parallel Execution Time :

$$T_p = t_{\text{comp}} + t_{\text{comm}} ; t_{\text{comp},i} = f_i(n, P) ;$$

$$t_{\text{comp}} = t_{\text{comp},1} + t_{\text{comp},2} + \dots ; \text{Similarly}$$

$$t_{\text{comm}} = t_{\text{comm},1} + t_{\text{comm},2} + \dots ; t_{\text{comm},i} = t_{\text{startup}} + w t_{\text{data}}$$

Lecture C: C-programming

* GDB

- gcc -g -Wall -o prog prog.c -lm
- gdb prog
 - . In gdb : type run to start your program
 - . bt : prints current call stack (list of nested functions)
 - . p x : prints value of variable x
 - . break file.c:123 : set break point
 - . continue : continues execution
 - . clear 1 : remove break point 1.
 - . l : list program code

* Valgrind

* Gprof

* IDE

Lecture 3: Introduction to MPI

Blocking send: `int MPI_Send`

Blocking Receive: `int MPI_Recv`

Non-blocking: `MPI_Isend`; `MPI_Wait`.

Data Movement Routines

- * Broadcast
- * Scattering \Rightarrow Gathering

Global Computation Routines

- * Reduce

MPI Calls:

- * `MPI_INIT`: establishes an environment.
- * `MPI_COMM_SIZE`: return # of processes
- * `MPI_COMM_RANK`: return rank of the process
- * `MPI_SEND`
- * `MPI_RECV`
- * `MPI_FINALIZE`: exit MPI cleanly.

Lecture 4. Image Reconstruction and Poisson's equation

The common denominator: use a uniform mesh.

* Load-balanced linear data distribution

* Load-balanced scatter distribution

Fill the Ghost Cells: Communication

\Rightarrow Safe solution: red-black coloring.

Red-Black Communication

- * Associate each process with a color (red or black), in the p and q directions such that no neighbor has the same color.

Optimal process Topology: Minimize $\Phi(P) = \left(\frac{P}{\mu} + \frac{Q}{N}\right)$

$$\Rightarrow P = \sqrt{\frac{\mu}{N}} R, R = P \times Q$$

Surface to Volume Ratio:

- * T_{comp} is proportional to $I_p * J_q$.
- * T_{comm} is proportional to $2(I_p + J_q)$

Gauss-Seidal Iteration: Change value immediately

Lexicographic Order: for $n=1..N$; for $m=1..M$: $u(m, n)$

Pipelined Computations: Gauss-Seidal. Red-black ordering.

Lecture 5. Matrix Multiplication and Collective Communication

Addition of Matrices:

Embarrassingly parallel, all data access purely local.

Recursive Doubling for taking sums.

Matrix-Vector Multiplication

- * Use a load-balanced data distribution on $R = P \times Q$ processors in array topology to store A .

- * Once data distributed, all individual components can be computed in parallel using the Recursive Doubling Algorithm

Matrix-Matrix Multiplication.

Assume $M=N=K$ and $R=P \times P$.

Cannon's Algorithm:

- * In a first step, assume $P=M$.
- * Data distribution: processor $(m-1, n-1)$ holds $a_{m,n}$, $b_{m,n}$, $c_{m,n}$
- * On the "diagonal" processors $(m-1, n-1)$, parts of the sum $c_{m,m} = a_{m,1} \cdot b_{1,m} + \dots + a_{m,m} \cdot b_{m,m} + \dots + a_{m,M} \cdot b_{M,m}$ are available.
- * Shifting the rows of A cyclic to the right and those of B cyclic downwards provides the next term.
- * Repeat this cyclic exchange ones again completes the calculation of the diagonal elements.

Algorithm:

- ① Initially, processor (p, q) has elements $a_{p+1, q+1}$, $b_{p+1, q+1}$
- ② Shift the rows of A and the columns of B into the structure described above.
- ③ For $k=1 \dots N$
 1. Multiply the local values on each processor
 2. Shift the rows of A and the columns of B cyclically by one processor
- ④ If necessary: undo step 1.

Definition of PageRank

- * Let m be a webpage
- * let F_m be the set of pages m points to
- * let B_m be the set of pages that point to m .
- * let $N_m = |F_m|$ be the number of forward links, and c be a normalization factor
- * Then we define PageRank $R_m = c \sum_{n \in B_m} \frac{R_n}{N_m}$
- * Define a matrix A by $a_{mn} = \begin{cases} 1/N_m, & \text{if there is link from } m \text{ to } n \\ 0, & \text{otherwise.} \end{cases}$
- * Let $R = (R_1, \dots, R_M)^T$ be the pagerank vector. Then it holds $R = cAR$.

The Power Method for Eigenvalue Problems.

- * Given a square matrix A and an initial guess $x^{[0]} \neq 0$
- * Form the sequence $x^{[k+1]} = Ax^{[k]}$, $k=0, 1, \dots$
- * For many A, $\{x^{[k]} / \|x^{[k]}\|\}$ converges towards an eigenvector to the largest eigenvalue c of A.
- * An estimation c_k of c is given by $c_k = \frac{\langle x^{[k+1]}, x^{[k]} \rangle}{\langle x^{[k]}, x^{[k]} \rangle}$
- * If $P=Q$, MPI-Alltoall can be used to transpose y .
- * The general case is equivalent to one recursive doubling step.

Lecture 6. Linear Systems of Equations by Gaussian Elimination

Gaussian Elimination

Idea: For $n=1, \dots, N-1$: Use equation n for the elimination of x_n from the equations $n+1, \dots, N$

For $n=N, \dots, 1$: Compute x_n from the n -th equation

Forward elimination + Backsubstitution

Pivoting + LU decomposition

Use data partition to distribute

Forward substitution: Sequential algorithm

Use pipelined computation: Lots of communication

Use Data Distribution μ :

- * Assume a row data distribution μ
- * The last process ($P-1$) must wait for x_1 to arrive.
- * x_i is ready to send if the first processor (0) has done its work.
- * In an ideal world, further communication is completely overlapped by computation.
- * $P \times Q$ processor mesh gives rather poor performance.

Lecture 7. Parallel Sorting

Symmetric Distributed Compare And Exchange

Linear data distribution place N/P consecutive elements on each processor. First step is local sorting.

Then do the merging. Complexity $(2N/P-1)t_a$.

Parallel Bubble Sort: Odd-Even Sort

2 phases, all processes are comparing. Guaranteed to terminate after $N/2$ odd-even and even-odd steps.

Bitonic Mergesort: Bitonic Sort implemented on

$P=N=2^D$ processor has a time complexity of $O(\log^3 N)$.

Bucket Sort: Uniformly distributed keys.

Lecture 8. Algorithms on Graph

Assume (connected) weighted undirected graph.

Algorithms For Dense Graphs

Prim's algorithm: Inner loop can be parallelized.

Distribute V load balanced over P processes.

Distribute the weighted adjacency matrix A correspondingly over a $1 \times P$ process mesh.

Single Source Shortest Paths - Dijkstra's Algorithm
Similar to Prim's algorithm

Floyd's Algorithm: Parallel innermost loop

Algorithms For Sparse Graphs